

RDR_ASSET Lineage

How CERD replication data becomes the published RDR_ASSET views — two landing tables reconciled into one view, two staging packages, and a ~20-step orchestration in between.


Table

View

Package

Procedure

External / unconfirmed



CERD data lands in two tables (**CERD_CORPORATE** primary, **CERD_CORP_CR** replica), gets reconciled by a UNION view, and is normalized through two staging packages — the second of which drives ~20 sequential procedures (p05→p180) — before loading **LRIDBO.RDR_ASSET**, the table three separate consumer views read from.

Open questions worth tracking

The upstream replication tool feeding CERD_CORPORATE and CERD_CORP_CR is not yet confirmed — it behaves like a GoldenGate-style CDC flow, but that's an assumption, not a verified fact. Anyone extending this pipeline should confirm the actual replication mechanism before relying on its latency or ordering guarantees.

STGCRDBO.LOAD_RDR_ASSET is the highest-risk hop: a single package orchestrating ~20 procedures in a fixed sequence (p05 through p180) means a failure anywhere in that chain blocks everything downstream of it, including all three published views.